
▼ Ungraded Lab: TensorFlow Model Analysis

In production systems, the decision to deploy a model usually goes beyond the global metrics (e.g. accuracy) set during training. It is also important to evaluate how your model performs in different scenarios. For instance, does your weather forecasting model perform equally well in summer compared to winter? Or does your camera-based defect detector work only in certain lighting conditions? This type of investigation helps to ensure that your model can handle different cases. More than that, it can help uncover any learned biases that can result in a negative experience for your users. For example, if you're supposed to have a gender-neutral application, you don't want your model to only work well for one while poorly for another.

In this lab, you will be working with [TensorFlow Model Analysis \(TFMA\)](#) – a library built specifically for analyzing a model's performance across different configurations. It allows you to specify slices of your data, then it will compute and visualize how your model performs on each slice. You can also set thresholds that your model must meet before it is marked ready for deployment. These help you make better decisions regarding any improvements you may want to make to boost your model's performance and ensure fairness.

For this exercise, you will use TFMA to analyze models trained on the [Census Income dataset](#). Specifically, you will:

- study and setup the starter files to use with TFMA
- make a configuration file to tell TFMA what data slices it will analyze and the metrics it will compute
- visualize TFMA's outputs in a notebook environment
- generate a time series of a model's performance
- compare the performance of two models so you can decide which one to push to production

Credits: Some of the code and discussions are based on the TensorFlow team's [official tutorial](#).

▼ Setup

In this section, you will first setup your workspace to have all the modules and files to work with TFMA. You will

- install required libraries,
- download starter files that will contain the dataset, schema, and pretrained models you will analyze
- prepare the dataset so it can be consumed by TFMA
- observe how the models transform the raw features

Install Jupyter Extensions

If running in a local Jupyter notebook, then these Jupyter extensions must be installed in the environment before running Jupyter. These are already available in Colab so we'll just leave the commands here for reference.

```
jupyter nbextension enable --py widgetsnbextension --sys-prefix
jupyter nbextension install --py --symlink tensorflow_model_analysis --sys-prefix
jupyter nbextension enable --py tensorflow_model_analysis --sys-prefix
```

▼ Install libraries

This will pull in all the dependencies, and will take a minute.

```
1 # Upgrade pip to the latest, and install required libraries.
2 !pip install -U pip
3 !pip install tensorflow_data_validation==1.1.0
4 !pip install tensorflow-transform==1.0.0
5 !pip install tensorflow-model-analysis==0.32.0
```

*Note: In Google Colab, you need to restart the runtime at this point to finalize updating the packages you just installed. **Please do not proceed to the next section without restarting.** You can also ignore the errors about version incompatibility of some of the bundled packages because we won't be using those in this notebook.*

▼ Check Installation

Running the code below should show the versions of the packages. Please re-run the install if you are seeing errors and don't forget to restart the runtime after re-installation.

```
1 # Import packages and print versions
2 import tensorflow as tf
3 import tensorflow_model_analysis as tfma
4 import tensorflow_data_validation as tfdv
5
6 print('TF version: {}'.format(tf.__version__))
7 print('TFMA version: {}'.format(tfma.__version__))
8 print('TFDV version: {}'.format(tfdv.__version__))
```

TF version: 2.5.2
TFMA version: 0.32.0
TFDV version: 1.1.0

▼ Load The Files

Next, you will download the files you will need for this exercise:

- Test datasets
- Data schema
- Pretrained models

We've also defined some global variables below so you can access these files throughout the notebook more easily.

```
1 import os
2
3 # String variables for file and directory names
4 URL = 'https://storage.googleapis.com/mlep-public/course_3/wee
5 TAR_NAME = 'C3_W4_Lab_1_starter_files.tar.gz'
6 BASE_DIR = 'starter_files'
7 DATA_DIR = os.path.join(BASE_DIR, 'data')
8 CSV_DIR = os.path.join(DATA_DIR, 'csv')
9 TFRECORD_DIR = os.path.join(DATA_DIR, 'tfrecord')
10 MODELS_DIR = os.path.join(BASE_DIR, 'models')
11 SCHEMA_FILE = os.path.join(BASE_DIR, 'schema.pbtxt')

1 # uncomment this line if you've downloaded the files before an
2 # !rm -rf {BASE_DIR}
3
4 # Download the tar file from GCP
5 !wget {URL}
6
7 # Extract the tar file to the base directory
8 !tar xzf {TAR_NAME}
9
10 # Delete tar file
11 !rm {TAR_NAME}
```

```
--2022-01-01 18:37:58-- https://storage.googleapis.com/mlep-public/course\_3/wee
Resolving storage.googleapis.com (storage.googleapis.com)... 172.253.122.128, 14
Connecting to storage.googleapis.com (storage.googleapis.com)|172.253.122.128|:4
HTTP request sent, awaiting response... 200 OK
```

```
Length: 1449046 (1.4M) [application/x-gzip]
Saving to: 'C3_W4_Lab_1_starter_files.tar.gz'
```

```
C3_W4_Lab_1_starter 100%[=====>] 1.38M --.-KB/s in 0.04s
```

```
2022-01-01 18:37:58 (32.4 MB/s) - 'C3_W4_Lab_1_starter_files.tar.gz' saved [1449
```

You can see the top level file and directories by running the cell below (or just using the file explorer on the left side of this Colab). We'll discuss what each contain in the next sections.

```
1 print("Here's what we downloaded:")
2 !ls {BASE_DIR}
```

```
Here's what we downloaded:
data models schema.pbtxt
```

▼ Preview the dataset

The `data/csv` directory contains the test split of the Census Income dataset. We've divided it into several files for this demo notebook:

- `data_test.csv` - 15000 rows of test data
- `data_test_1.csv` - first 5000 rows of `data_test.csv`
- `data_test_2.csv` - next 5000 rows of `data_test.csv`
- `data_test_3.csv` - last 5000 rows of `data_test.csv`

You can see the description of each column [here](#) (please open link in a new window if Colab prevents the download). Also for simplicity, we've already preprocessed the `label` column as binary (i.e. 0 or 1) to match the model's output. In your own projects, your labels might be in a different data type (e.g. string) and you want to transform that first so you can evaluate your model properly. You can preview the first few rows below:

```
1 # Path to the full test set
2 TEST_DATA_PATH = os.path.join(CSV_DIR, 'data_test.csv')
3
4 # Preview the first few rows
5 !head {TEST_DATA_PATH}
```

```
age,workclass,fnlwgt,education,education-num,marital-status,occupation,relations
25,Private,226802,11th,7,Never-married,Machine-op-inspct,Own-child,Black,Male,0,
38,Private,89814,HS-grad,9,Married-civ-spouse,Farming-fishing,Husband,White,Male
28,Local-gov,336951,Assoc-acdm,12,Married-civ-spouse,Protective-serv,Husband,Whi
44,Private,160323,Some-college,10,Married-civ-spouse,Machine-op-inspct,Husband,E
```

```
18,?,103497,Some-college,10,Never-married,?,Own-child,White,Female,0,0,30,United
34,Private,198693,10th,6,Never-married,Other-service,Not-in-family,White,Male,0,
29,?,227026,HS-grad,9,Never-married,?,Unmarried,Black,Male,0,0,40,United-States,
63,Self-emp-not-inc,104626,Prof-school,15,Married-civ-spouse,Prof-specialty,Husb
24,Private,369667,Some-college,10,Never-married,Other-service,Unmarried,White,Fe
```

▼ Parse the Schema

You also downloaded a schema generated by [TensorFlow Data Validation](#). You should be familiar with this file type already from previous courses. You will load it now so you can use it in the later parts of the notebook.

```
1 import pandas as pd
2 d = pd.read_csv(TEST_DATA_PATH)
3 d.head()
```

1 to 5 of 5 entries Filter ?

index	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	r
0	25	Private	226802	11th	7	Never-married	Machine-op-inspct	Own-child	B
1	38	Private	89814	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband	V
2	28	Local-gov	336951	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husband	V
3	44	Private	160323	Some-college	10	Married-civ-spouse	Machine-op-inspct	Husband	B
4	18	?	103497	Some-college	10	Never-married	?	Own-child	V

Show 25 ▼ per page


Like what you see? Visit the [data table notebook](#) to learn more about interactive tables.

```
1 d.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 16281 entries, 0 to 16280
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                    16281 non-null  int64
1   workclass              16281 non-null  object
2   fnlwgt                 16281 non-null  int64
3   education              16281 non-null  object
4   education-num          16281 non-null  int64
5   marital-status         16281 non-null  object
6   occupation             16281 non-null  object
7   relationship           16281 non-null  object
8   race                   16281 non-null  object
9   sex                    16281 non-null  object
```

```
10 capital-gain      16281 non-null  int64
11 capital-loss      16281 non-null  int64
12 hours-per-week    16281 non-null  int64
13 native-country    16281 non-null  object
14 label             16281 non-null  int64
dtypes: int64(7), object(8)
memory usage: 1.9+ MB
```

```
1 # Load the schema as a protocol buffer
2 SCHEMA = tfdv.load_schema_text(SCHEMA_FILE)
3
4 # Display the schema
5 tfdv.display_schema(SCHEMA)
```

	Type	Presence	Valency	Domain	
Feature name					
'age'	INT	required		-	
'capital-gain'	INT	required		-	
'capital-loss'	INT	required		-	
'education'	STRING	required		'education'	
'education-num'	INT	required		-	
'fnlwgt'	INT	required		-	
'hours-per-week'	INT	required		-	
'label'	INT	required		-	
'marital-status'	STRING	required		'marital-status'	

▼ Use the Schema to Create TFRecords

TFMA needs a TFRecord file input so you need to convert the CSVs in the data directory. If you've done the earlier labs, you will know that this can be easily done with `ExampleGen`. For this notebook however, you will use the helper function below instead to demonstrate how it can be done outside a TFX pipeline. You will pass in the schema you loaded in the previous step to determine the correct type of each feature.

```

1 # imports for helper function
2 import csv
3 from tensorflow.core.example import example_pb2
4 from tensorflow_metadata.proto.v0 import schema_pb2
5
6 def csv_to_tfrecord(schema, csv_file, tfrecord_file):
7     ''' Converts a csv file into a tfrecord
8     Args:
9         schema (schema_pb2) - Schema protobuf from TFDV
10        csv_file (string) - file to convert to tfrecord
11        tfrecord_file (string) - filename of tfrecord to create
12
13    Returns:
14        filename of tfrecord
15    '''
16
17    # Open CSV file for reading. Each row is mapped as a dictio
18    reader = csv.DictReader(open(csv_file, 'r'))

```

```

19
20 # Initialize TF examples list
21 examples = []
22
23 # For each row in CSV, create a TF Example based on
24 # the Schema and append to the list
25 for line in reader:
26
27     # Intialize example
28     example = example_pb2.Example()
29
30     # Loop through features in the schema
31     for feature in schema.feature:
32
33         # Get current feature name
34         key = feature.name
35
36         # Populate values based on data type of current feature
37         if feature.type == schema_pb2.FLOAT:
38             example.features.feature[key].float_list.value[:] = (
39                 [float(line[key])] if len(line[key]) > 0 else [])
40         elif feature.type == schema_pb2.INT:
41             example.features.feature[key].int64_list.value[:] = (
42                 [int(line[key])] if len(line[key]) > 0 else [])
43         elif feature.type == schema_pb2.BYTES:
44             example.features.feature[key].bytes_list.value[:] = (
45                 [line[key].encode('utf8')] if len(line[key]) > 0 else [])
46
47     # Append to the list
48     examples.append(example)
49
50 # Write examples to tfrecord file
51 with tf.io.TFRecordWriter(tfrecord_file) as writer:
52     for example in examples:
53         writer.write(example.SerializeToString())
54
55 return tfrecord_file

```

The code below will do the conversion and we've defined some more global variables that you will use in later exercises.

```

1 # Create tfrecord directory

```



```

2 !mkdir {TFRECORD_DIR}
3
4 # Create list of tfrecord files
5 tfrecord_files = [csv_to_tfrecord(SCHEMA, f'{CSV_DIR}/{name}',
6   for name in os.listdir(CSV_DIR)]
7
8 # Print created files
9 print(f'files created: {tfrecord_files}')
10
11 # Create variables for each tfrecord
12 TFRECORD_FULL = os.path.join(TFRECORD_DIR, 'data_test.tfrecord
13 TFRECORD_DAY1 = os.path.join(TFRECORD_DIR, 'data_test_1.tfreco
14 TFRECORD_DAY2 = os.path.join(TFRECORD_DIR, 'data_test_2.tfreco
15 TFRECORD_DAY3 = os.path.join(TFRECORD_DIR, 'data_test_3.tfreco
16
17 # Delete unneeded variable
18 del tfrecord_files

files created: ['starter_files/data/tfrecord/data_test.tfrecord', 'starter_files

```

▼ Pretrained models

Lastly, you also downloaded pretrained Keras models and they are stored in the `models/` directory. TFMA supports a number of different model types including TF Keras models, models based on generic TF2 signature APIs, as well TF estimator based models. The [get_started](#) guide has the full list of model types supported and any restrictions. You can also consult the [FAQ](#) for examples on how to configure these models.

We have included three models and you can choose to analyze any one of them in the later sections. These were saved in [SavedModel](#) format which is the default when saving with the Keras Models API.

```

1 # list model directories
2 !ls {MODELS_DIR}
3
4 # Create string variables for each model directory
5 MODEL1_FILE = os.path.join(MODELS_DIR, 'model1')
6 MODEL2_FILE = os.path.join(MODELS_DIR, 'model2')
7 MODEL3_FILE = os.path.join(MODELS_DIR, 'model3')

model1  model2  model3

```

As mentioned earlier, these models were trained on the [Census Income dataset](#). The label is 1 if a person earns more than 50k USD and 0 if less than or equal. You can load one of the models and look at the summary to get a sense of its architecture. All three models use the same architecture but were trained with different epochs to simulate varying performance.

```
1 # Load model 1
2 model = tf.keras.models.load_model(MODEL1_FILE)
3
4 # Print summary. You can ignore the warnings at the start.
5 model.summary()
```

```
WARNING:tensorflow:SavedModel saved prior to TF 2.5 detected when loading Keras
WARNING:tensorflow:SavedModel saved prior to TF 2.5 detected when loading Keras
WARNING:tensorflow:Inconsistent references when loading the checkpoint into this
```

```
Two checkpoint references resolved to different objects (<tensorflow.python.kera
WARNING:tensorflow:Inconsistent references when loading the checkpoint into this
```

```
Two checkpoint references resolved to different objects (<tensorflow.python.kera
Model: "model"
```

Layer (type)	Output Shape	Param #	Connected to
fnlwgt_xf (InputLayer)	[(None, 1)]	0	
education-num_xf (InputLayer)	[(None, 1)]	0	
capital-gain_xf (InputLayer)	[(None, 1)]	0	
capital-loss_xf (InputLayer)	[(None, 1)]	0	
hours-per-week_xf (InputLayer)	[(None, 1)]	0	
concatenate (Concatenate)	(None, 5)	0	fnlwgt_xf[0][0] education-num_x capital-gain_xf capital-loss_xf hours-per-week_
dense (Dense)	(None, 100)	600	concatenate[0][
dense_1 (Dense)	(None, 70)	7070	dense[0][0]
education_xf (InputLayer)	[(None, 21)]	0	
marital-status_xf (InputLayer)	[(None, 12)]	0	
occupation_xf (InputLayer)	[(None, 20)]	0	
race_xf (InputLayer)	[(None, 10)]	0	
relationship_xf (InputLayer)	[(None, 11)]	0	

workclass_xf (InputLayer)	[(None, 14)]	0	
sex_xf (InputLayer)	[(None, 7)]	0	
native-country_xf (InputLayer)	[(None, 47)]	0	
age_xf (InputLayer)	[(None, 4)]	0	
dense_2 (Dense)	(None, 48)	3408	dense_1[0][0]
concatenate_1 (Concatenate)	(None, 146)	0	education_xf[0] marital-status_ occupation_xf[0] race_xf[0][0] relationship_xf workclass_xf[0]

You can see the code to build these in the next lab. For now, you'll only need to take note of a few things. First, the output is a single dense unit with a sigmoid activation (i.e. `dense_5` above). This is standard for binary classification problems.

Another is that the model is exported with a transformation layer. You can see this in the summary above at the bottom row named `transform_features_layer` and it is not connected to the other layers. From previous labs, you will know that this is taken from the graph generated by the `Transform` component. It helps to avoid training-serving skews by making sure that raw inputs are transformed in the same way that the model expects. It is also available as a `tft_layer` property of the model object.

```
1 # Transformation layer can be accessed in two ways. These are
2 model.get_layer('transform_features_layer') is model.tft_layer

True
```

TFMA invokes this layer automatically for your raw inputs but we've included a short snippet below to demonstrate how the transformation works. You can see that the raw features are indeed reformatted to an acceptable input for the model. The raw numeric features are scaled and the raw categorical (string) features are encoded to one-hot vectors.

```
1 from tensorflow_transform.tf_metadata import schema_utils
2
3 # Load one tfrecord
4 tfrecord_file = tf.data.TFRecordDataset(TFRECORD_DAY1)
5
6 # Parse schema object as a feature spec
```



```
marital-status_xf: [[0. 1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]]  
occupation_xf: [[0. 0. 0. 0. 0. 0. 1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]]  
native-country_xf: [[1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.  
    0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]]  
capital-loss_xf: [0.]  
capital-gain_xf: [0.]  
fnlwgt_xf: [0.14569008]  
hours-per-week_xf: [0.39795917]
```

The transformed features can be passed into the model to get the predictions. The snippet below demonstrates this and we used a low-threshold [BinaryAccuracy](#) metric to compare the true labels and model predictions.

```

1 from tensorflow_transform.tf_metadata import schema_utils
2
3 # Load one tfrecord
4 tfrecord_file = tf.data.TFRecordDataset(TFRECORD_DAY1)
5
6 # Parse schema object as a feature spec
7 feature_spec = schema_utils.schema_as_feature_spec(SCHEMA).fea
8
9 # Create a batch from the dataset
10 for records in tfrecord_file.batch(5).take(1):
11
12     # Get the label values from the raw input
13     parsed_examples = tf.io.parse_example(records, feature_spec)
14     y_true = parsed_examples.pop('label')
15     print(f'labels:\n {y_true.numpy()}\n')
16
17     # Transform the raw features and pass to the model to get pr
18     transformed_examples = model.tft_layer(parsed_examples)
19     y_pred = model(transformed_examples)
20     print(f'predictions:\n {y_pred.numpy()}\n')
21
22     # Measure the binary accuracy
23     metric = tf.keras.metrics.BinaryAccuracy(threshold=0.3)
24     metric.update_state(y_true, y_pred)
25     print(f'binary accuracy: {metric.result().numpy()}\n')

```

labels:

```

[[0]
[0]
[1]
[1]

```

```
[0]]

predictions:
[[1.6402992e-34]
 [3.4708142e-02]
 [5.1936507e-03]
 [3.3919078e-01]
 [2.3632433e-15]]

binary accuracy: 0.800000011920929
```

Last thing to note is the model is also exported with a [serving signature](#). You will know more about this in the next lab and in later parts of the specialization but for now, you can think of it as a configuration for when the model is deployed for inference. For this particular model, the default signature is configured to transform batches of serialized raw features before feeding to the model input. That way, you wouldn't have to explicitly code the transformations as previously shown in the snippet above. You can just pass in batches of data directly as shown below.

```
1 # Load one tfrecord
2 tfrecord_file = tf.data.TFRecordDataset(TFRECORD_DAY1)
3
4 # Print available signatures
5 print(f'model signatures: {model.signatures}\n')
6
7 # Create a batch
8 for records in tfrecord_file.batch(5).take(1):
9
10     # Pass the batch to the model serving signature to get predictions
11     output = model.signatures['serving_default'](examples=record_batch)
12
13     # Print results
14     print(f"predictions:\n {output['output_0']}\n")

model signatures: _SignatureMap({'serving_default': <ConcreteFunction signature_
...
predictions:
[[1.6402867e-34]
 [3.4708112e-02]
 [5.1937401e-03]
 [3.3919084e-01]
 [2.3632433e-15]]
```

TFMA accesses this model signature so it can work with the raw data and evaluate the model to get the metrics. Not only that, it can also extract specific features and domain values from your

dataset before it computes these metrics. Let's see how this is done in the next section.

▼ Setup and Run TFMA

With the dataset and model now available, you can now move on to use TFMA. There are some additional steps needed:

- Create a `tfma.EvalConfig` protocol message containing details about the models, metrics, and data slices you'd like to analyze
- Create a `tfma.EvalSharedModel` that points to your saved models.
- Specify an output path where the analysis results will be stored

▼ Create EvalConfig

The `tfma.EvalConfig()` is a protocol message that sets up the analysis. Here, you will specify:

- [model_specs](#) - message containing at least the label key so it can be extracted from the evaluation/test data
- [metrics_specs](#) - containing the metrics you would like to evaluate. A comprehensive guide can be found [here](#) and you will use the binary classification metrics for this exercise.
- [slicing_specs](#) - containing the feature slices you would like to compute metrics for. A short description of different types of slices is shown [here](#)

The eval config should be passed as a protocol message and you can use the [google.protobuf.text_format](#) module for that as shown below.

```
1  import tensorflow_model_analysis as tfma
2  from google.protobuf import text_format
3
4  # Setup tfma.EvalConfig settings
5  eval_config = text_format.Parse("""
6      ## Model information
7      model_specs {
8          # For keras (and serving models), you need to add a `label
9          label_key: "label"
10     }
11
12     ## Post training metric information. These will be merged wi
13     ## metrics from training.
14     metrics_specs {
15         metrics { class_name: "ExampleCount" }
16         metrics { class_name: "BinaryAccuracy" }
```

```

16     metrics { class_name: "BinaryAccuracy" }
17     metrics { class_name: "BinaryCrossentropy" }
18     metrics { class_name: "AUC" }
19     metrics { class_name: "AUCPrecisionRecall" }
20     metrics { class_name: "Precision" }
21     metrics { class_name: "Recall" }
22     metrics { class_name: "MeanLabel" }
23     metrics { class_name: "MeanPrediction" }
24     metrics { class_name: "Calibration" }
25     metrics { class_name: "CalibrationPlot" }
26     metrics { class_name: "ConfusionMatrixPlot" }
27     # ... add additional metrics and plots ...
28 }
29
30 ## Slicing information
31
32 # overall slice
33 slicing_specs {}
34
35 # slice specific features
36 slicing_specs {
37     feature_keys: ["sex"]
38 }
39 slicing_specs {
40     feature_keys: ["race"]
41 }
42
43 # slice specific values from features
44 slicing_specs {
45     feature_values: {
46         key: "native-country"
47         value: "Cambodia"
48     }
49 }
50 slicing_specs {
51     feature_values: {
52         key: "native-country"
53         value: "Canada"
54     }
55 }
56
57 # slice feature crosses
58 slicing_specs {
59     - - - - -

```



```

59     feature_keys: ["sex", "race"]
60     }
61     "", tfma.EvalConfig())

```

▼ Create EvalSharedModel

TFMA also requires an [EvalSharedModel](#) instance that points to your model so it can be shared between multiple threads in the same process. This instance includes information about the type of model (keras, etc) and how to load and configure the model from its saved location on disk (e.g. tags, etc). The [tfma.default_eval_shared_model\(\)](#) API can be used to create this given the model location and eval config.

```

1 # Create a tfma.EvalSharedModel that points to your model.
2 # You can ignore the warnings generated.
3 eval_shared_model = tfma.default_eval_shared_model(
4     eval_saved_model_path=MODEL1_FILE,
5     eval_config=eval_config)

```

```

WARNING:tensorflow:SavedModel saved prior to TF 2.5 detected when loading Keras
WARNING:tensorflow:SavedModel saved prior to TF 2.5 detected when loading Keras
WARNING:tensorflow:Inconsistent references when loading the checkpoint into this

```

```

Two checkpoint references resolved to different objects (<tensorflow.python.kera
WARNING:tensorflow:Inconsistent references when loading the checkpoint into this

```

```

Two checkpoint references resolved to different objects (<tensorflow.python.kera

```

```

1 # Show properties of EvalSharedModel
2 print(f'EvalSharedModel contents: {eval_shared_model}')

```

```

EvalSharedModel contents: EvalSharedModel(model_path='starter_files/models/model

```

▼ Run TFMA

With the setup complete, you just need to declare an output directory then run TFMA. You will pass in the eval config, shared model, dataset, and output directory to [tfma.run_model_analysis\(\)](#) as shown below. This will create a [tfma.EvalResult](#) which you can use later for rendering metrics and plots.

```

1 # Specify output path for the evaluation results
2 OUTPUT_DIR = os.path.join(BASE_DIR, 'output')
3

```

```
4 # Run TFMA. You can ignore the warnings generated.
5 eval_result = tfma.run_model_analysis(
6     eval_shared_model=eval_shared_model,
7     eval_config=eval_config,
8     data_location=TFRECORD_FULL,
9     output_path=OUTPUT_PATH)
```

```
WARNING:absl:Tensorflow version (2.5.2) found. Note that TFMA support for TF 2.0
WARNING:apache_beam.runners.interactive.interactive_environment:Dependencies req
WARNING:tensorflow:SavedModel saved prior to TF 2.5 detected when loading Keras
WARNING:tensorflow:SavedModel saved prior to TF 2.5 detected when loading Keras
WARNING:tensorflow:Inconsistent references when loading the checkpoint into this
```

```
Two checkpoint references resolved to different objects (<tensorflow.python.kera
WARNING:tensorflow:Inconsistent references when loading the checkpoint into this
```

```
Two checkpoint references resolved to different objects (<tensorflow.python.kera
WARNING:root:Make sure that locally built Python SDK docker image has Python 3.7
WARNING:tensorflow:SavedModel saved prior to TF 2.5 detected when loading Keras
WARNING:tensorflow:SavedModel saved prior to TF 2.5 detected when loading Keras
WARNING:tensorflow:Inconsistent references when loading the checkpoint into this
```

```
Two checkpoint references resolved to different objects (<tensorflow.python.kera
WARNING:tensorflow:Inconsistent references when loading the checkpoint into this
```

```
Two checkpoint references resolved to different objects (<tensorflow.python.kera
WARNING:apache_beam.io.tfrecordio:Couldn't find python-snappy so the implementat
WARNING:tensorflow:SavedModel saved prior to TF 2.5 detected when loading Keras
```

▼ Visualizing Metrics and Plots

You can visualize the results also using TFMA methods. In this section, you will view the returned metrics and plots for the different slices you specified in the eval config.

```
WARNING:tensorflow:SavedModel saved prior to TF 2.5 detected when loading Keras
```

▼ Rendering Metrics

You can view the metrics with the [tfma.view.render_slicing_metrics\(\)](#) method. By default, the views will display the `overall` slice. To view a particular slice you can pass in a feature name to the `slicing_column` argument as shown below. You can visualize the different metrics through the `show` dropdown menu and you can hover over the bar charts to show the exact value measured.

We encourage you to try the different options you see and also modify the command. Here are some examples:

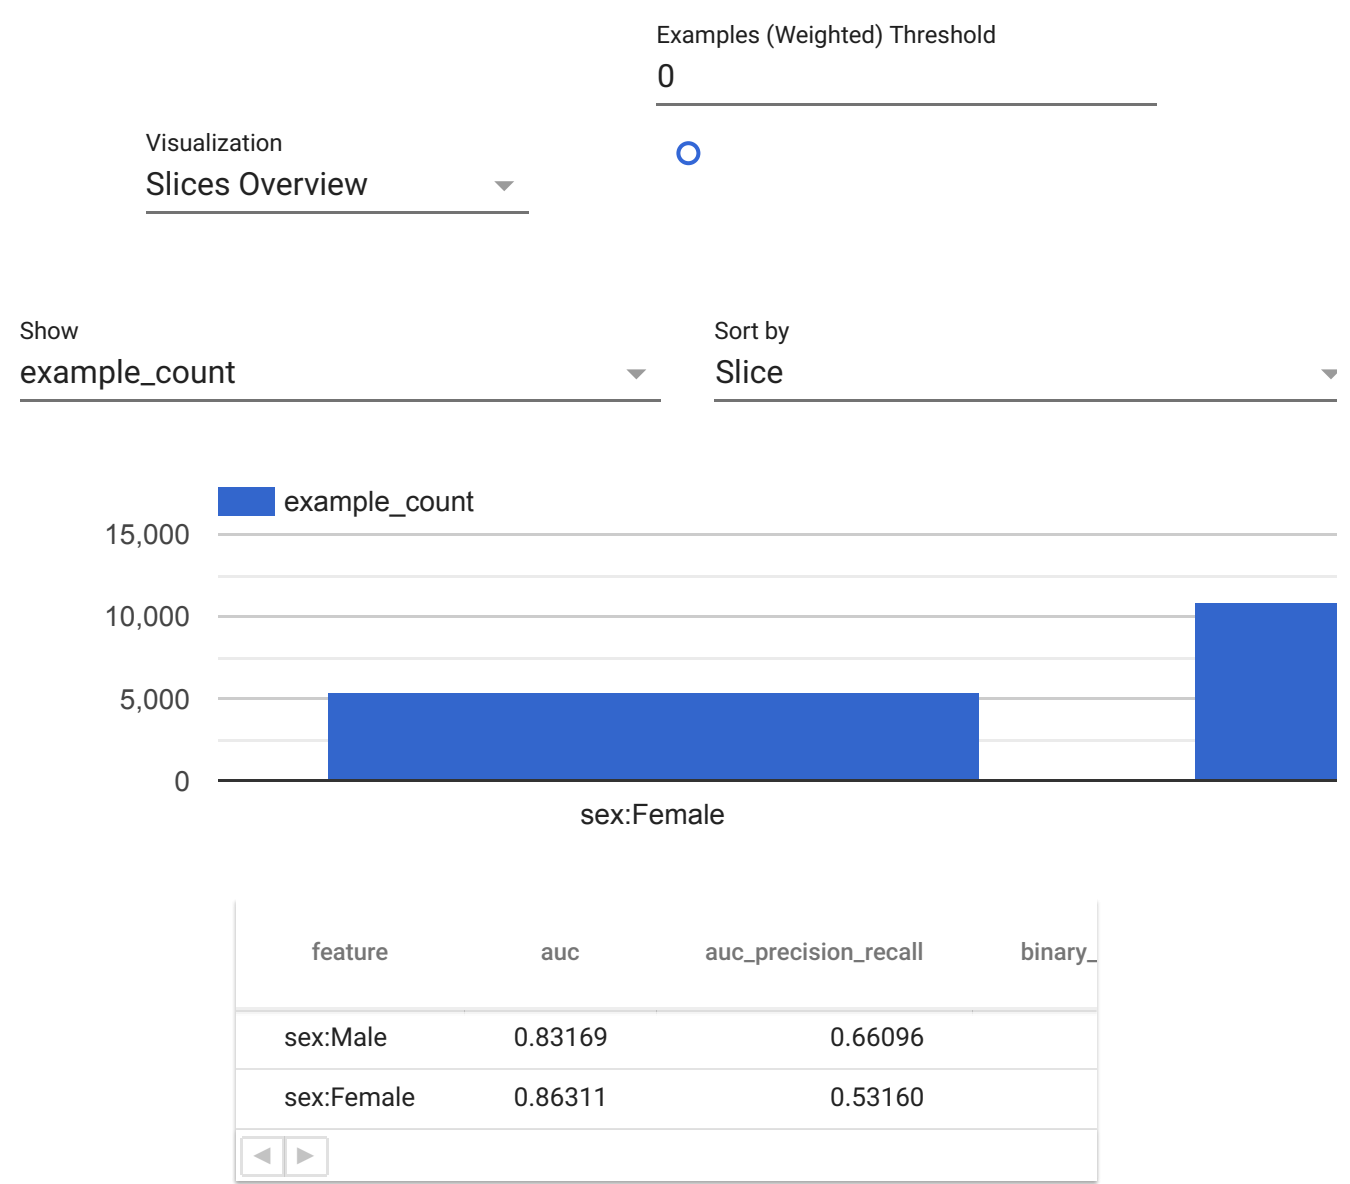
- Removing the `slicing_column` argument will produce the overall slice.
- You can also pass in `race` (since it was specified in the eval config) to see the results for that particular slice.
- Using the `Examples (Weighted) Threshold` slider above 5421 will remove the `Female` slice because it has less examples than that.
- Toggling the `view` dropdown to `Metrics Histogram` will show the results divided into buckets. For example, if you're slicing column is `sex` and the `Histogram Type` dropdown is at `slice Counts`, then you will one slice in two of the 10 (default) buckets since we only

have two values for that feature ('Male' and 'Female'). The x-axis show the values for the metric in the `Select Metric` dropdown. This is the default view when the number of slices is large.

- At the bottom of the screen, you will notice the measurements also presented in tabular format. You can sort it by clicking on the feature name headers.

```
> tf.data.Dataset.from_tensor_slices([0, 1])
```

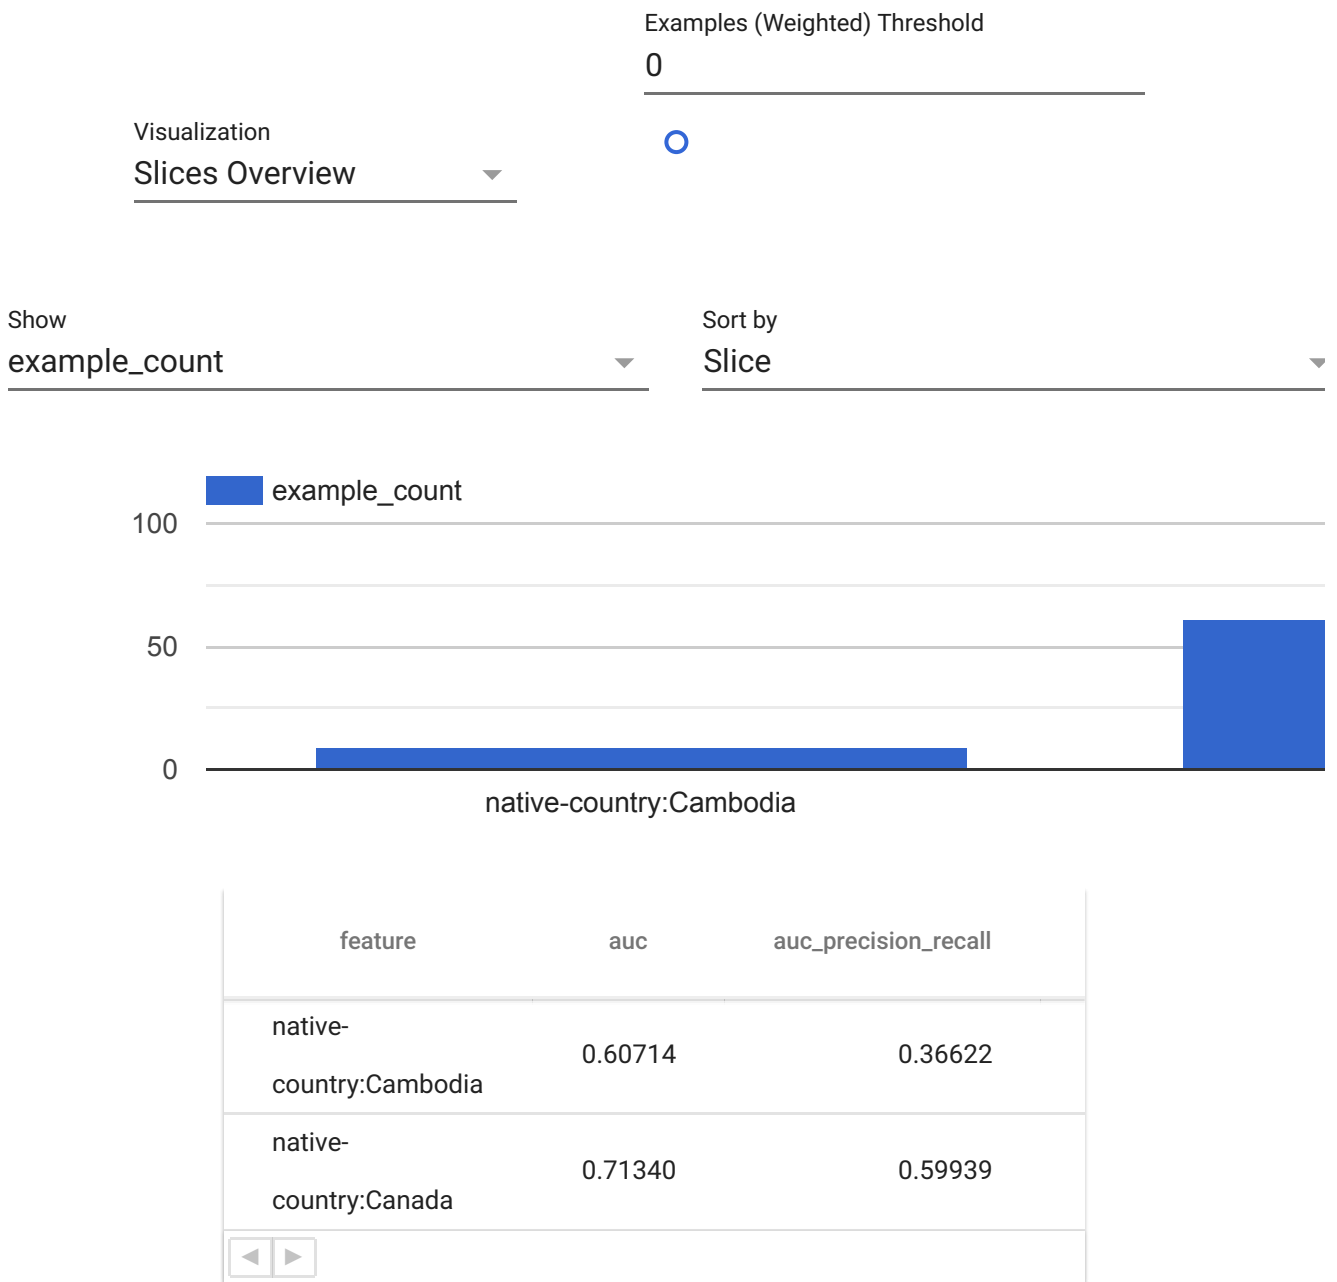
```
1 # Render metrics for a feature
2 tfma.view.render_slicing_metrics(eval_result, slicing_column='sex')
```



▼ More Slices

If you haven't yet, you can also pass in the `native-country` to the slicing column. The difference in this visualization is we only specified two of its values in the eval config earlier. This is useful if you just want to study a subgroup of a particular feature and not the entire domain.

```
1 # Render metrics for feature. Review EvalConfig message to see
2 tfma.view.render_slicing_metrics(eval_result, slicing_column='native-country')
```



TFMA also supports creating feature crosses to analyze combinations of features. Our original settings created a cross between `sex` and `race` and you can pass it in as a [SlicingSpec](#) as shown below.

```
1 # Render metrics for feature crosses
2 tfma.view.render_slicing_metrics(
3     eval_result,
4     slicing_spec=tfma.SlicingSpec(
5         feature_keys=['sex', 'race']))
```

Examples (Weighted) Threshold

0

Visualization

Slices Overview



Show

example count

Sort by

Slice

In some cases, crossing the two columns creates a lot of combinations. You can narrow down the results to only look at specific values by specifying it in the `slicing_spec` argument. Below shows the results for the `sex` feature for the `Other` race.

```
1 # Narrow down the feature crosses by specifying feature values
2 tfma.view.render_slicing_metrics(
3     eval_result,
4     slicing_spec=tfma.SlicingSpec(
5         feature_keys=['sex'], feature_values={'race': 'Other'})
```

Examples (Weighted) Threshold

0

Visualization

Slices Overview



Show

Sort by

▼ Rendering Plots

Any plots that were added to the `tfma.EvalConfig` as post training `metric_specs` can be displayed using [tfma.view.render_plot](#).

As with metrics, plots can be viewed by slice. Unlike metrics, only plots for a particular slice value can be displayed so the `tfma.SlicingSpec` must be used and it must specify both a slice feature name and value. If no slice is provided then the plots for the `overall` slice is used.

The example below displays the plots that were computed for the `sex:Male` slice. You can click on the names at the bottom of the graph to see a different plot type. Alternatively, you can tick the `Show all plots` checkbox to show all the plots in one screen.

```
1 # Render plots
2 tfma.view.render_plot(
3     eval_result,
4     tfma.SlicingSpec(feature_values={'sex': 'Male'}))
```


Plot for sex:Male



▼ Tracking Model Performance Over Time

Your training dataset will be used for training your model, and will hopefully be representative of your test dataset and the data that will be sent to your model in production. However, while the data in inference requests may remain the same as your training data, it can also start to change enough so that the performance of your model will change. That means that you need to monitor and measure your model's performance on an ongoing basis so that you can be aware of and react to changes.

Let's take a look at how TFMA can help. You will load three different datasets and compare the model analysis results using the [render_time_series\(\)](#) method.

 Show all plots

```
1  import os
2
3  # Put data paths we prepared earlier in a list
4  TFRECORDS = [TFRECORD_DAY1, TFRECORD_DAY2, TFRECORD_DAY3]
5
6  # Initialize output paths list for each result
7  output_paths = []
8
9  # Run eval on each tfrecord separately
10 for num, tfrecord in enumerate(TFRECORDS):
11
12     # Use the same model as before
13     eval_shared_model = tfma.default_eval_shared_model(
14         eval_saved_model_path=MODEL1_FILE,
15         eval_config=eval_config)
16
17     # Prepare output path name
18     output_path = os.path.join('.', 'time_series', str(num))
19     output_paths.append(output_path)
20
21     # Run TFMA on the current tfrecord in the loop
22     tfma.run_model_analysis(eval_shared_model=eval_shared_model,
23                             eval_config=eval_config,
24                             data_location=tfrecord
```

Two checkpoint references resolved to different objects (<tensorflow.python.kera
 WARNING:absl:Tensorflow version (2.5.2) found. Note that TFMA support for TF 2.0
 WARNING:tensorflow:SavedModel saved prior to TF 2.5 detected when loading Keras
 WARNING:tensorflow:SavedModel saved prior to TF 2.5 detected when loading Keras
 WARNING:tensorflow:Inconsistent references when loading the checkpoint into this

Two checkpoint references resolved to different objects (<tensorflow.python.kera
 WARNING:tensorflow:Inconsistent references when loading the checkpoint into this

Two checkpoint references resolved to different objects (<tensorflow.python.kera
 WARNING:root:Make sure that locally built Python SDK docker image has Python 3.7
 WARNING:tensorflow:SavedModel saved prior to TF 2.5 detected when loading Keras
 WARNING:tensorflow:SavedModel saved prior to TF 2.5 detected when loading Keras
 WARNING:tensorflow:Inconsistent references when loading the checkpoint into this

Two checkpoint references resolved to different objects (<tensorflow.python.kera
 WARNING:tensorflow:Inconsistent references when loading the checkpoint into this

Two checkpoint references resolved to different objects (<tensorflow.python.kera
 WARNING:tensorflow:SavedModel saved prior to TF 2.5 detected when loading Keras
 WARNING:tensorflow:SavedModel saved prior to TF 2.5 detected when loading Keras
 WARNING:tensorflow:Inconsistent references when loading the checkpoint into this

Two checkpoint references resolved to different objects (<tensorflow.python.kera
 WARNING:tensorflow:Inconsistent references when loading the checkpoint into this

Two checkpoint references resolved to different objects (<tensorflow.python.kera
 WARNING:tensorflow:SavedModel saved prior to TF 2.5 detected when loading Keras
 WARNING:tensorflow:SavedModel saved prior to TF 2.5 detected when loading Keras
 WARNING:tensorflow:Inconsistent references when loading the checkpoint into this

Two checkpoint references resolved to different objects (<tensorflow.python.kera
 WARNING:tensorflow:Inconsistent references when loading the checkpoint into this

Two checkpoint references resolved to different objects (<tensorflow.python.kera
 WARNING:tensorflow:SavedModel saved prior to TF 2.5 detected when loading Keras
 WARNING:tensorflow:SavedModel saved prior to TF 2.5 detected when loading Keras
 WARNING:tensorflow:Inconsistent references when loading the checkpoint into this

Two checkpoint references resolved to different objects (<tensorflow.python.kera
 WARNING:tensorflow:Inconsistent references when loading the checkpoint into this

Two checkpoint references resolved to different objects (<tensorflow.python.kera
 WARNING:tensorflow:SavedModel saved prior to TF 2.5 detected when loading Keras
 WARNING:tensorflow:SavedModel saved prior to TF 2.5 detected when loading Keras
 WARNING:tensorflow:Inconsistent references when loading the checkpoint into this

Two checkpoint references resolved to different objects (<tensorflow.python.kera
 WARNING:tensorflow:Inconsistent references when loading the checkpoint into this

Two checkpoint references resolved to different objects (<tensorflow.python.kera
 WARNING:tensorflow:SavedModel saved prior to TF 2.5 detected when loading Keras
 WARNING:tensorflow:SavedModel saved prior to TF 2.5 detected when loading Keras
 WARNING:tensorflow:Inconsistent references when loading the checkpoint into this

```
WARNING:tensorflow:Inconsistent references when loading the checkpoint into this
```

```
Two checkpoint references resolved to different objects (<tensorflow.python.kera
```

```
WARNING:tensorflow:Inconsistent references when loading the checkpoint into this
```

First, imagine that you've trained and deployed your model yesterday. And now, you want to see how it's doing on the new data coming in today. The visualization will start by displaying AUC. From the UI, you can:

- Add other metrics using the "Add metric series" menu.
- Close unwanted graphs by clicking on x
- Hover over data points (the ends of line segments in the graph) to get more details

Note: In the metric series charts, the x-axis is just the model directory name of the model that you're examining.

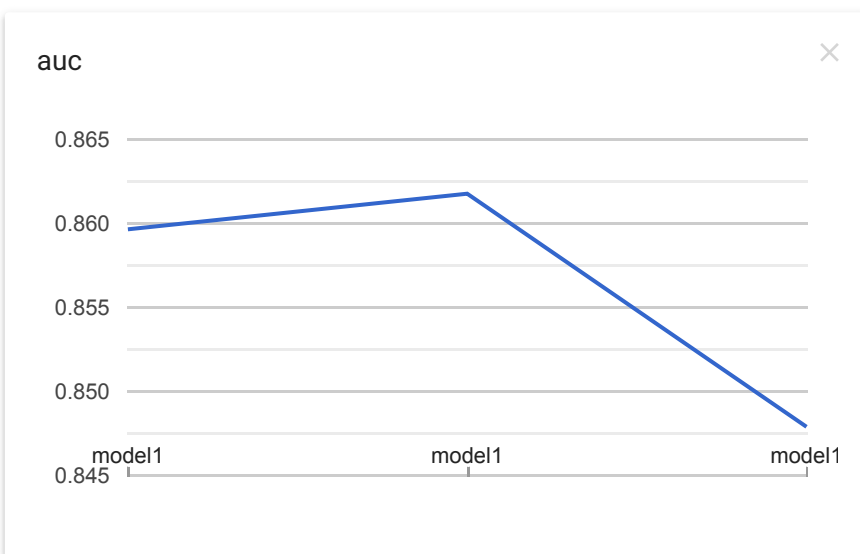
```
1 # Load results for day 1 and day 2 datasets
2 eval_results_from_disk = tfma.load_eval_results(output_paths[:
3
4 # Visualize results
5 tfma.view.render_time_series(eval_results_from_disk)
```

Add metric series

Now imagine that another day has passed and you want to see how it's doing on the new data coming in today.

```
1 # Load results for all three days
2 eval_results_from_disk = tfma.load_eval_results(output_paths)
3
4 # Visualize the results
5 tfma.view.render_time_series(eval_results_from_disk)
```

Add metric series



Model	Data	auc	auc_precision_recall	binary_accuracy
model1	data_test_3.tfrecord	0.84791	0.63176	0.82697
model1	data_test_2.tfrecord	0.86177	0.65618	0.82837
model1	data_test_1.tfrecord	0.85965	0.64573	0.82900

This type of investigation lets you see if your model is behaving poorly on new data. You can make the decision to retrain your production model based on these results. Retraining might not always produce the best results and you also need a way to detect that. You will see how TFMA helps in that regard in the next section.

▼ Model Validation

TFMA can be configured to evaluate multiple models at the same time. Typically, this is done to compare a candidate model against a baseline (such as the currently serving model) to determine what the performance differences in metrics are. When [thresholds](#) are configured, TFMA will produce a [tfma.ValidationResult](#) record indicating whether the performance matches expectations.

Below, you will re-configure the EvalConfig settings to compare two models: a candidate and a baseline. You will also validate the candidate's performance against the baseline by setting a [tmfa.MetricThreshold](#) on the BinaryAccuracy metric. This helps in determining if your new model can indeed replace your currently deployed model.

```
1  # Setup tfma.EvalConfig setting with metric thresholds
2  eval_config_with_thresholds = text_format.Parse("""
3      ## Model information
4      model_specs {
5          name: "candidate"
6          label_key: "label"
7      }
8      model_specs {
9          name: "baseline"
10         label_key: "label"
11         is_baseline: true
12     }
13
14     ## Post training metric information
15     metrics_specs {
16         metrics { class_name: "ExampleCount" }
17         metrics {
18             class_name: "BinaryAccuracy"
19             threshold {
20                 # Ensure that metric is always > 0.9
21                 value_threshold {
22                     lower_bound { value: 0.9 }
23                 }
24                 # Ensure that metric does not drop by more than a small
25                 # e.g. (candidate - baseline) > -1e-10 or candidate >
26                 change_threshold {
27                     direction: HIGHER_IS_BETTER
28                     absolute { value: -1e-10 }
29                 }
30             }
```

```

31     }
32     metrics { class_name: "BinaryCrossentropy" }
33     metrics { class_name: "AUC" }
34     metrics { class_name: "AUCPrecisionRecall" }
35     metrics { class_name: "Precision" }
36     metrics { class_name: "Recall" }
37     metrics { class_name: "MeanLabel" }
38     metrics { class_name: "MeanPrediction" }
39     metrics { class_name: "Calibration" }
40     metrics { class_name: "CalibrationPlot" }
41     metrics { class_name: "ConfusionMatrixPlot" }
42     # ... add additional metrics and plots ...
43 }
44
45 ## Slicing information
46 slicing_specs {} # overall slice
47 slicing_specs {
48     feature_keys: ["race"]
49 }
50 slicing_specs {
51     feature_keys: ["sex"]
52 }
53 """, tfma.EvalConfig())
54
55 # Create tfma.EvalSharedModels that points to the candidate and
56 candidate_model_path = MODEL1_FILE
57 baseline_model_path = MODEL2_FILE
58
59 eval_shared_models = [
60     tfma.default_eval_shared_model(
61         model_name=tfma.CANDIDATE_KEY,
62         eval_saved_model_path=candidate_model_path,
63         eval_config=eval_config_with_thresholds),
64     tfma.default_eval_shared_model(
65         model_name=tfma.BASELINE_KEY,
66         eval_saved_model_path=baseline_model_path,
67         eval_config=eval_config_with_thresholds),
68 ]
69
70 # Specify validation path
71 validation_output_path = os.path.join(OUTPUT_DIR, 'validation')
72
73 # Run TFMA on the two models

```

```

13 # RUN TFMA ON THE TWO MODELS
74 eval_result_with_validation = tfma.run_model_analysis(
75     eval_shared_models,
76     eval_config=eval_config_with_thresholds,
77     data_location=TFRECORD_FULL,
78     output_path=validation_output_path)

```

Two checkpoint references resolved to different objects (<tensorflow.python.keras.SavedModel object at 0x7f8b1c1b1b1b>
WARNING:tensorflow:SavedModel saved prior to TF 2.5 detected when loading Keras
WARNING:tensorflow:SavedModel saved prior to TF 2.5 detected when loading Keras
WARNING:tensorflow:Inconsistent references when loading the checkpoint into this

Two checkpoint references resolved to different objects (<tensorflow.python.keras.SavedModel object at 0x7f8b1c1b1b1b>
WARNING:tensorflow:Inconsistent references when loading the checkpoint into this

Two checkpoint references resolved to different objects (<tensorflow.python.keras.SavedModel object at 0x7f8b1c1b1b1b>
WARNING:tensorflow:SavedModel saved prior to TF 2.5 detected when loading Keras
WARNING:tensorflow:SavedModel saved prior to TF 2.5 detected when loading Keras
WARNING:tensorflow:Inconsistent references when loading the checkpoint into this

Two checkpoint references resolved to different objects (<tensorflow.python.keras.SavedModel object at 0x7f8b1c1b1b1b>
WARNING:tensorflow:Inconsistent references when loading the checkpoint into this

Two checkpoint references resolved to different objects (<tensorflow.python.keras.SavedModel object at 0x7f8b1c1b1b1b>
WARNING:tensorflow:SavedModel saved prior to TF 2.5 detected when loading Keras
WARNING:tensorflow:SavedModel saved prior to TF 2.5 detected when loading Keras
WARNING:tensorflow:Inconsistent references when loading the checkpoint into this

Two checkpoint references resolved to different objects (<tensorflow.python.keras.SavedModel object at 0x7f8b1c1b1b1b>
WARNING:tensorflow:Inconsistent references when loading the checkpoint into this

Two checkpoint references resolved to different objects (<tensorflow.python.keras.SavedModel object at 0x7f8b1c1b1b1b>
WARNING:tensorflow:SavedModel saved prior to TF 2.5 detected when loading Keras
WARNING:tensorflow:SavedModel saved prior to TF 2.5 detected when loading Keras
WARNING:tensorflow:Inconsistent references when loading the checkpoint into this

Two checkpoint references resolved to different objects (<tensorflow.python.keras.SavedModel object at 0x7f8b1c1b1b1b>
WARNING:tensorflow:Inconsistent references when loading the checkpoint into this

Two checkpoint references resolved to different objects (<tensorflow.python.keras.SavedModel object at 0x7f8b1c1b1b1b>
WARNING:tensorflow:SavedModel saved prior to TF 2.5 detected when loading Keras
WARNING:tensorflow:SavedModel saved prior to TF 2.5 detected when loading Keras
WARNING:tensorflow:Inconsistent references when loading the checkpoint into this

Two checkpoint references resolved to different objects (<tensorflow.python.keras.SavedModel object at 0x7f8b1c1b1b1b>
WARNING:tensorflow:Inconsistent references when loading the checkpoint into this

Two checkpoint references resolved to different objects (<tensorflow.python.keras.SavedModel object at 0x7f8b1c1b1b1b>
WARNING:tensorflow:SavedModel saved prior to TF 2.5 detected when loading Keras
WARNING:tensorflow:SavedModel saved prior to TF 2.5 detected when loading Keras
WARNING:tensorflow:Inconsistent references when loading the checkpoint into this

Two checkpoint references resolved to different objects (<tensorflow.python.keras.SavedModel object at 0x7f8b1c1b1b1b>
WARNING:tensorflow:Inconsistent references when loading the checkpoint into this

Two checkpoint references resolved to different objects (<tensorflow.python.kera
WARNING:tensorflow:SavedModel saved prior to TF 2.5 detected when loading Keras
WARNING:tensorflow:SavedModel saved prior to TF 2.5 detected when loading Keras
WARNING:tensorflow:Inconsistent references when loading the checkpoint into this

Two checkpoint references resolved to different objects (<tensorflow.python.kera
WARNING:tensorflow:Inconsistent references when loading the checkpoint into this


Two checkpoint references resolved to different objects (<tensorflow.python.kera

When running evaluations with one or more models against a baseline, TFMA automatically adds different metrics for all the metrics computed during the evaluation. These metrics are named after the corresponding metric but with the string `_diff` appended to the metric name. A positive value for these `_diff` metrics indicates an improved performance against the baseline.

Like in the previous section, you can view the results with `render_time_series()`.

```
1 # Render results
2 tfma.view.render_time_series(eval_result_with_validation)
```

Add metric series



Model	Data	auc	auc_diff	auc_precision_recall	auc_pre
model2	data_test.tfrecord	0.85297	n/a	0.61562	
model1	data_test.tfrecord	0.85684	3.8697e-3	0.64060	

You can use [tfma.load_validator_result](#) to view the validation results you specified with the threshold settings. For this example, the validation fails because `BinaryAccuracy` is below the

```
1 # Print validation result
2 validation_result = tfma.load_validation_result(validation_out
3 print(validation_result.validation_ok)
```

```
False
```

Congratulations! You have now explored the different methods of model analysis using TFMA. In the next section, you will see how these can fit into a TFX pipeline so you can automate the process and store the results in your pipeline directory and metadata store.